

## Assignment -3

### Abalone Age Prediction

Assignment Date	10 October 2022
Student Name	TAMIL PRABHA R
Student Roll Number	820419104077
Maximum Marks	2 Marks

#### 1.Importing necessary packages & Downloading the packages

```
import pandas as pd
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt import seaborn as
sns
from sklearn.preprocessing import StandardScaler from
sklearn.preprocessing import LabelEncoder import numpy
as np from collections import Counter from
sklearn.pipeline import make_pipeline from
sklearn.linear_model import Ridge, Lasso from
sklearn.ensemble import RandomForestRegressor from
sklearn.pipeline import make_pipeline from
sklearn.linear_model import Ridge, Lasso from
sklearn.model_selection import GridSearchCV from
sklearn.exceptions import NotFittedError from
sklearn.metrics import r2_score,mean_absolute_error
```

#### 2. Download the dataset:

```
df= pd.read_csv("abalone.csv")
```

```
df.head()
```

#### Output:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15


	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

### 3. Visualizations (i) Univariate Analysis

**Input:** sns.displot(df["Diameter"])

**Output:**


<seaborn.axisgrid.FacetGrid at 0x1a7c3cc60a0>



**Input:**

```
sns.set(rc={'figure.figsize':(7,7)})
depth = df['Rings'].value_counts(normalize=False).reset_index()
sns.barplot(data=depth,x='index',y='Rings') Output:
```


<AxesSubplot:xlabel='index', ylabel='Rings'>



### Input:

`df['Sex'].value_counts().plot(kind='pie',autopct='%0f')` **Output:**

<AxesSubplot:ylabel='Sex'>




### (ii) BiVariate Analysis

#### Input:

`sns.set(rc={'figure.figsize':(17,10)})`

`sns.countplot(df['Rings'],hue = df['Sex'],color = 'y')` **Output:**

<AxesSubplot:xlabel='Rings', ylabel='count'>

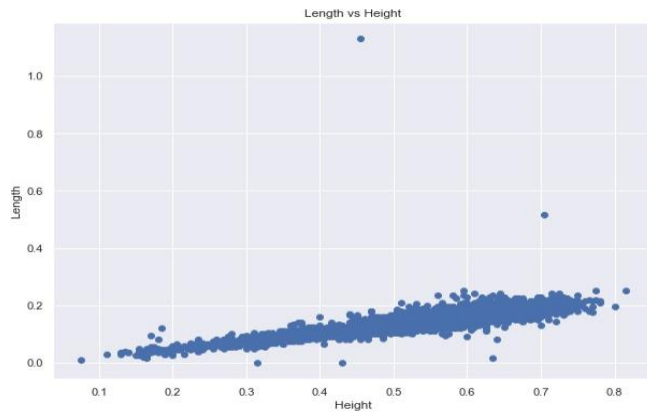


### Input:

```
sns.set(rc={'figure.figsize':(10,7)})
plt.scatter(df.Length, df.Height)
plt.title('Length vs Height')
plt.xlabel('Height') plt.ylabel('Length')
```

### Output:

Text(0, 0.5, 'Length')



### (iii) MultiVariate Analysis


#### Input:

```
plt.figure(figsize=(12,10))
sns.pairplot(df)
```

#### Output:

<seaborn.axisgrid.PairGrid at 0x1a8005d43a0>

<Figure size 864x720 with 0 Axes>




### Input:

```
plt.figure(figsize = (8,6)) corr =
df.corr() sns.heatmap(corr,
annot = True)
```

### Output:

<AxesSubplot:>



#### 4.Descriptive Statistics

**Input:** df.info()

**Output:**

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 4177 entries, 0 to 4176

Data columns (total 9 columns):

# Column Non-Null Count Dtype

0 Sex 4177 non-null object

1 Length 4177 non-null float64

2 Diameter 4177 non-null float64

3 Height 4177 non-null float64

4 Whole weight 4177 non-null float64

5 Shucked weight 4177 non-null float64

6 Viscera weight 4177 non-null float64

7 Shell weight 4177 non-null float64 8 Rings 4177 non-null int64 dtypes: float64(7), int64(1), object(1) memory usage: 293.8+ KB

**Input:**

df.describe()

**Output:**

	<b>Length</b>	<b>Diameter</b>	<b>Height</b>	<b>Whole weight</b>	<b>Shucked weight</b>	<b>Viscera weight</b>	<b>Shell weight</b>	<b>Rings</b>
<b>count</b>	4177.00000	4177.00000	4177.00000	4177.00000	4177.00000	4177.00000	4177.00000	4177.00000
<b>mean</b>	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.238831	9.933684
<b>std</b>	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.139203	3.224169
<b>min</b>	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001500	1.000000
<b>25%</b>	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.130000	8.000000
<b>50%</b>	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.234000	9.000000
<b>75%</b>	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.329000	11.000000
<b>max</b>	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000	1.005000	29.000000



	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
x	0	0	0	0	0	0	0	00

## 5.Handle Missing Values

### Input:

```
df.isna().sum()
```

### Output:

```
Sex          0
Length       0
Diameter     0
Height       0
Whole weight 0
Shucked weight 0
Viscera weight 0
Shell weight 0
Rings        0
```

```
dtype: int64 6.
```

### Outlier Detection

**Input:** outlier\_correction\_df =  
df.drop(columns=['Sex'],axis=1)  
outlier\_correction\_df.head()

### Output:

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15

<b>1</b>	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
	<b>Length</b>	<b>Diameter</b>	<b>Height</b>	<b>Whole weight</b>	<b>Shucked weight</b>	<b>Viscera weight</b>	<b>Shell weight</b>	<b>Rings</b>
<b>2</b>	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
<b>3</b>	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
<b>4</b>	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

### Input:

**def** detection(df,features):

outlier\_indices=[]

**for** c **in** features:

    Q1 = np.percentile(df[c],25)

    Q3 = np.percentile(df[c],75)

    IQR = Q3 - Q1      outlier\_step  
= IQR \* 1.5      lower\_range = Q1 -  
(outlier\_step)

    upper\_range = Q3 + (outlier\_step)

    outlier\_list\_col=df[ (df[c] < lower\_range) | (df[c] > upper\_range) ].index

    outlier\_indices.extend(outlier\_list\_col)

**return** outlier\_indices

**def** multiple\_outlier\_indices(outlier\_indices):

```

outlier_indices=Counter(outlier_indices)
multiple_outliers = list(i for i, v in outlier_indices.items() if v > 2 )

return multiple_outliers

```

**Solution:** outlier\_correction\_df.columns

### Output:

```

Index(['Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight',
      'Viscera weight', 'Shell weight', 'Rings'],
      dtype='object')

```

### Solution:

```

outliers=detection(df,['Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight',
      'Viscera weight', 'Shell weight'])

```

**Input:** Counter(outliers)

### Output:

```

Counter({ 148: 2,

```

```

      149: 2,
      236: 3,
      237: 3,
      238: 3,
      239: 3,
      305: 2,
      306: 3,
      321: 2,
      465: 2,
      523: 2,
      525: 2,
      526: 2,
      611: 2,
      694: 3,
      696: 2,
      718: 3,
      719: 3,
      720: 3,

```

1053: 2,  
1054: 2,  
1055: 2,  
1056: 1,  
1210: 1,  
1429: 3,  
1824: 2,  
1986: 2,  
1987: 3,  
2114: 3,  
2115: 2,  
2169: 3,  
2171: 3,  
2343: 2,  
2371: 2,  
2380: 2,  
2381: 3,  
2458: 2,  
2711: 3,  
3141: 2,  
3143: 2,  
3190: 3,  
3318: 2,  
3380: 2,  
3472: 2,  
3600: 2,  
3837: 3,  
3899: 3,  
3902: 3,  
3994: 2,  
43: 1,  
44: 1,  
520: 1,  
892: 1,  
898: 1,  
1988: 1,  
2172: 2,  
2545: 1,  
2712: 1,  
3473: 1,  
3521: 1,  
3716: 1,  
1174: 1,

1257: 1,  
1417: 2,  
1428: 3,  
1763: 4,  
2051: 1,  
2179: 1,  
3996: 1,  
165: 3,  
358: 2,  
891: 3,  
1051: 2,  
1052: 3,  
1193: 3,  
1206: 3, 1207: 4,  
1209: 3,  
1426: 2,  
1427: 3,  
1761: 3,  
1762: 4,  
2265: 1,  
2334: 2,  
2623: 3,  
2624: 3,  
2811: 3,  
2862: 2,  
2863: 3,  
3007: 2,  
3008: 2,  
3188: 2,  
3427: 3,  
3599: 2,  
3715: 4,  
3800: 1,  
3993: 2,  
1048: 2,  
1197: 1,  
1199: 1,  
1202: 1,  
1418: 1,  
1527: 1,  
1528: 1,  
1749: 1,  
1750: 2,

1754: 1,  
1756: 1,  
1821: 1,  
1982: 1,  
2544: 1,  
2625: 1,  
2675: 1,  
2710: 2,  
2810: 2,  
2970: 1,  
2972: 1,  
3082: 1,  
3713: 1,  
3961: 1,  
3962: 1,  
170: 1, 1204: 1,  
1422: 1,  
1757: 1,  
1759: 1,  
2709: 1,  
3628: 1,  
4148: 1,  
81: 1,  
129: 1,  
157: 1,  
163: 1,  
164: 1,  
166: 1,  
167: 1,  
168: 1,  
277: 1,  
334: 1,  
1823: 1,  
1985: 1,  
2090: 1,  
2108: 1,  
2157: 1,  
2161: 1,  
2208: 1,  
2274: 1,  
2368: 1,  
3148: 1,  
3149: 1,

```
3151: 1,  
3928: 1,  
4145: 1})
```

**Input:**

```
multiple_outlier_indices = multiple_outlier_indices(outliers)
```

**Solution:** `print(Counter(multiple_outlier_indices))`

```
Counter({236: 1, 237: 1, 238: 1, 239: 1, 306: 1, 694: 1, 718: 1, 719: 1, 720: 1, 1429: 1, 1987: 1,  
2114: 1, 2169: 1, 2171: 1, 2381: 1, 2711: 1, 3190: 1, 3837: 1, 3899: 1, 3902: 1, 1428: 1, 1763: 1,  
165: 1, 891: 1, 1052: 1, 1193: 1, 1206: 1, 1207: 1, 1209: 1, 1427: 1, 1761: 1, 1762: 1, 2623: 1, 2624:  
1, 2811: 1, 2863: 1, 3427: 1, 3715: 1})
```

**Input:** `df=df.drop(multiple_outlier_indices,axis=0).reset_index(drop`  
**= True) Output:**

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1500	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.0700	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.2100	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.1550	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.0550	7
...	...	...	...	...	...	...	...	...	...
4134	F	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	11
4135	M	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605	10
4136	M	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080	9
4137	F	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960	10
4138	M	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.4950	12

4139 rows × 9 columns

**Input:**

df.shape

**Output:**

(4139, 9)

## 7. Categorical Attribute Encoding

**Input:** le=LabelEncoder()

df['Sex']=le.fit\_transform(df['Sex'])

) **Solution:** df.head() **Output:**



	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	2	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

### 8. Seperate dataframe into Predictor and Target

#### Input:

```
feature = pd.DataFrame(df.drop(['Rings'], axis = 1))
label = pd.DataFrame(df.Rings)
```

### 9. Scaling the Predictor variables Input:

```
convert = StandardScaler()
feature = pd.DataFrame(convert.fit_transform(feature))
```

### 10. Perform the train test split

#### Input:

```
X_train, X_test, y_train, y_test = train_test_split(feature, label, test_size = 0.2, random_state = 0)
```

#### Input:

```
print('X_train : ')
print(X_train)
print(X_train.shape)
```

```
print("") print('X_test
: ') print(X_test)
print(X_test.shape)
```

```
print("") print('y_train
: ') print(y_train)
print(y_train.shape)
```

```
print("")
print('y_test : ')
print(y_test)
print(y_test.shape)
X_train :
```

```
      0      1      2      3      4      5      6 \
64  1.151942 -0.040971 -0.087769 -0.481759 -0.514480 -0.574405 -0.453057
1521 -0.062807 -1.409929 -1.432469 -1.091508 -1.310493 -1.322352 -1.158735
3436 1.151942 1.670227 1.618965 1.469438 2.013909 1.674082 1.916340
3444 -0.062807 -0.768230 -0.760119 -0.725659 -0.868496 -0.720742 -0.798886
3993 -0.062807 -0.725450 -0.863557 -0.725659 -0.689393 -0.462910 -0.565218
...      ...      ...      ...      ...      ...      ...
1033 1.151942 1.413547 1.205212 0.859689 1.832711 2.173488 1.701365
3264 1.151942 1.028527 0.946616 1.225538 0.845026 0.772829 0.916240
1653 -1.277555 0.729068 0.688019 0.493839 0.607270 0.120118 0.402169
2607 -0.062807 -0.725450 -0.708400 -0.969558 -0.876875 -0.690546 -0.995167
2732 -0.062807 0.215709 0.119108 0.371889 0.180984 0.638106 0.004933
```

```
      7
64  -0.390700
1521 -1.300351
3436 2.132846
3444 -1.014251
3993 -0.959232
...      ...
1033 1.223195
3264 1.149837
1653 1.032462
2607 -0.992243
2732 -0.243983
```

```
[3311 rows x 8 columns]
(3311, 8)
```

```
X_test :
      0      1      2      3      4      5      6 \
```

```

958  1.151942 -0.126531  0.015669  0.127990 -0.062009  0.134054  0.014280
2613 -0.062807 -0.425991 -0.346365 -0.603709 -0.639119 -0.579050 -0.569891
45   -0.062807 -1.153250 -1.173873 -1.091508 -1.304208 -1.254990 -1.261549
3145 -1.277555 -0.169311  0.015669 -0.115910 -0.353182 -0.309604 -0.439037  3994
-0.062807 -0.340431 -0.449804 -1.213458 -0.365751 -0.191140 -0.448384
...   ...   ...   ...   ...   ...   ...   ...
620  -1.277555 -0.853790 -0.760119 -1.091508 -1.054931 -1.101684 -1.112002
1544 -0.062807 -0.597110 -0.501523 -0.603709 -0.772137 -0.692869 -0.883007
2954  1.151942  0.087369 -0.036050  0.859689  0.931959  0.884324  1.369556
177  -0.062807 -2.564988 -2.570292 -2.311006 -1.632040 -1.545343 -1.541951
50   -0.062807 -0.040971  0.015669 -0.481759 -0.483058 -0.553499 -0.644665

```

```

      7
958  -0.313674
2613 -0.680468
45   -1.197648
3145 -0.317342
3994 -0.427380
...   ...
620  -0.830854
1544 -0.669464
2954  0.724355
177  -1.637802
50   -0.354021

```

[828 rows x 8 columns]  
(828, 8)

```

y_train :
Rings
64      8
1521    8
3436   11
3444    7
3993    8
...   ...
1033    8
3264   17
1653   10
2607    7
2732    9

```

[3311 rows x 1 columns]

(3311, 1)

y\_test :

Rings

958 8

2613 7

45 7

3145 15

3994 8

... ...

620 10

1544 10

2954 13

177 4

50 8

[828 rows x 1 columns]

(828, 1)

## 11.Build Model

**Input:** pipelines={

'rf':make\_pipeline(RandomForestRegressor(random\_state=1234)),

'ridge':make\_pipeline(Ridge(random\_state=1234)),

'lasso':make\_pipeline(Lasso(random\_state=1234)),

} **Input:** hyperparagrid={

'rf':{

'randomforestregressor\_\_min\_samples\_split':[2,4,6],

'randomforestregressor\_\_min\_samples\_leaf':[1,2,3]

},

'ridge':{

'ridge\_\_alpha':[0.001,0.005,0.01,0.05,0.1,0.5,0.99]

},

'lasso':{

'lasso\_\_alpha':[0.001,0.005,0.01,0.05,0.1,0.5,0.99]

}

}

## 12. Training the Model

### Input:

```
fit_models={} for algo, pipeline in
pipelines.items():
    model=GridSearchCV(pipeline,hyperparagrid[algo],cv=10,n_jobs=-1)
try:
    print('Start training for {}'.format(algo))
    model.fit(X_train,y_train)
    fit_models[algo]=model except
NotFittedError as e:    print(repr(e)) Start
training for rf
Start training for ridge
Start training for lasso
```

## 13,14 Testing and Measuring Performance

**Input:** best\_model\_rf=fit\_models['rf']

best\_model\_rf **Output:**

```
GridSearchCV(cv=10,
    estimator=Pipeline(steps=[('randomforestregressor',
        RandomForestRegressor(random_state=1234))]),
n_jobs=-1,
    param_grid={'randomforestregressor__min_samples_leaf': [1, 2, 3],
        'randomforestregressor__min_samples_split': [2, 4, 6]})
```

### Input:

best\_model\_ridge=fit\_models['ridge']

best\_model\_ridge **Output:**

```
GridSearchCV(cv=10,
    estimator=Pipeline(steps=[('ridge', Ridge(random_state=1234))]),
n_jobs=-1,
    param_grid={'ridge__alpha': [0.001, 0.005, 0.01, 0.05, 0.1, 0.5,
        0.99]})
```

**Input:** best\_model\_lasso=fit\_models['lasso']

best\_model\_lasso

**Output:**

```
GridSearchCV(cv=10,
             estimator=Pipeline(steps=[('lasso', Lasso(random_state=1234))]),
             n_jobs=-1,
             param_grid={'lasso__alpha': [0.001, 0.005, 0.01, 0.05, 0.1, 0.5,
                                           0.99]}))
```

**Input:**

```
for algo,model in fit_models.items():
    ya=model.predict(X_test)
    print('{ } scores-R2:{ } MAE:{ }'.format(algo,r2_score(y_test,ya),
    mean_absolute_error(y_test,ya))) Output:
```

```
rf    scores-R2:0.5255029479701915    MAE:1.570513566816263
ridge scores-R2:0.5189099860811324 MAE:1.6528099660919895
lasso scores-R2:0.5190720174119673 MAE:1.6525494856846143
```